

Dmitry Roitman 150 West End Avenue, Brooklyn, NY 11235 | Phone: 718 404 6471 | dmitryro@gmail.com

LinkedIn: <https://www.linkedin.com/in/dmitry-r-4104b56/> | Website: <https://linktr.ee/dmitryro>

| GitHub: <https://github.com/dmitryro>

OBJECTIVE: With over 20 years of experience in application development and architecture, I am seeking senior, staff-level or principal-level roles that leverage my expertise while offering opportunities for growth and knowledge sharing.

TECHNICAL SKILLS:

Programming Languages:

- Python: Flask, Django, Django Rest Framework, FastAPI, Tornado, Pandas, Polars, Numpy, Pydantic, SQLAlchemy, Alembic, Pydantic (15 years)
- Java: Spring, Spring Boot, Spring Rest, Play, Websphere, JBoss, Tomcat, Akka (20 years)
- Go: Web Application, Middleware, System Development, GORM, Gin, Chi, Mux (12 years)
- C# and .NET: .NET Core, ASP.NET (8 years)
- Scala: AKKA (5 years)
- Rust: System Development, Middleware, Serde, async-std, Tokio, Rocket, Actix Web (5 years)
- JavaScript / TypeScript: Angular, React/Redux/Next.js, Vue/Nuxt.js, Node.js, Express.js, RXJS, D3.js, Three.js, Svelte, Tailwind CSS, SASS/SCSS (12 years)

Cloud and DevOps:

- Cloud Infrastructure: AWS (IAM, EC2, S3, VPC, CloudFormation, Lambda, ELB, SageMaker, AWS Textract, boto3, etc.) (13 years), Google Cloud, GC Document AI, MS Azure (10 years)
- CI/CD: Jenkins, Docker, Kubernetes, ArgoCD, Terraform, Sentry, Pulumi, Ansible, Looker, Datadog, New Relic, Prometheus, Grafana, OpenTelemetry, GitLab, Github Actions, ELK Stack, MLflow, Kubeflow, DVC, Dagster, Prefect (12 years)

Data Engineering and Messaging:

- Data Engineering: Apache Airflow, Luigi, OpenLineage, Snowflake, BigQuery, Redshift, DBT, Databricks, Apache Spark, Cassandra, Parquet, Apache Arrow, Apache Flink (12 years).
- Messaging: Celery, Apache Kafka, AWS Kinesis, RabbitMQ, AMQP, MQTT, WebSocket, Samza, ZeroMQ, Debezium (12 years)

Machine Learning and AI:

- Custom Model Development, Training and Deployment: NumPy, SciPy, PyTorch, PyJanitor, Pandas, Keras, Tensorflow, Fast.ai, Spacy, Stanza, Gensim, NLTK, Scikit-Learn (12 years).
- Third Party LLM and RAG: OpenAI, Hugging Face, LangChain, LlamaIndex (5 years)

Databases:

- RDBMS/SQL: PostgreSQL, MySQL, Cosmos DB, CockroachDB, SQL (25 years)
- NoSQL: Redis, CouchDB, DynamoDB, Neo4J, Cypher, Memgraph, ElasticSearch (15 years)

Integration, Security, Networking and Compliance:

- Integration: REST, GraphQL, gRPC, SOAP, FHIR/HL7/X12/EDI, HAPI-FHIR, JWT (15 years)
- Security: SSL/TLS, OpenSSL, RSA, LDAP, OAuth2, JWT, Okta, SAML, OpenID (15 years)
- Networking: TCP, UDP, MQTT, AMQP, HTTP, HTTPS, Websocket (15 years)
- Compliance: SoC2, HIPAA, PCI (10 years)

Linux System Administration and Virtualization:

- Sys. Administration: Shell Scripting, Debian/Ubuntu, RHEL, Bash, AWK, Backups (15 years)
- Virtualization: ProxMox, Qemu, VMWare, Virtuozzo, Xen (15 years)

Data Scraping, Search and Visualization:

- Data Scraping: BeautifulSoup, Scrapy (10 years)
- Visualizations: D3.js, Three.js, HighCharts, Tableau, PowerBI, OpenGL (15 years)
- Apache Solr, Haystack, Lucene, Xapian, ElasticSearch (10 years)

Version Control:

- Git (GitHub, Bitbucket) (15 years), Subversion (10 years)

Software Architecture, Leadership and Software Development Lifecycle:

- SDLC: TOGAF, Six Sigma, Zachman, Solutions Architecture, Agile (20 years)
- System Design: Domain Driven Design, Event Driven Architectures, Applications of SOLID/ DRY, CQRS, Architecture Patterns, Microservices, Tradeoff Analysis (15 years)
- Soft Skills: Technical Leadership, Mentorship, Stakeholder Management, Knowledge Transfer, Team Building (20 years).

Dmitry Roitman 150 West End Avenue, Brooklyn, NY 11235 | Phone: 718 404 6471 | dmitryro@gmail.com

LinkedIn: <https://www.linkedin.com/in/dmitry-r-4104b56/> | Website: <https://linktr.ee/dmitryro>

| GitHub: <https://github.com/dmitryro>

LANGUAGES:

- English (full fluency)
- Russian (full fluency)
- Hebrew (full fluency)

EDUCATION:

From: 09/1995. To: 08/1999.

UNIVERSITY OF HAIFA, Haifa, Israel

Degree Earned: Bachelor of Science in Math and Computer Science

CERTIFICATIONS: SUN CERTIFIED JAVA PROGRAMMER

PROFESSIONAL EXPERIENCE:

Senior Staff Engineer

Feb 2021 – Now

Legend Energy Advisors, New York, NY

- Architected, implemented and documented the corporate analytics SaaS platform used to monitor and optimize utility (electricity, gas, water) consumption and customer budgets, leveraging advanced analytics and actionable insights.
- Architected and built scalable session management, data streaming, and analytics dashboards with Python, PostgreSQL, RabbitMQ, FastAPI, Django, Angular, and WebSocket.
- Developed and maintained RabbitMQ-based messaging infrastructure to enable seamless communication and scaling across the platform.
- Designed and integrated dynamic data visualization tools (dashboards, charts, analysis) using D3.js, Three.js, HighCharts, and Pandas.
- Architected and built observability initiatives with Go, Kafka, Debezium, ELK Stack, PostgreSQL, and Angular for enhanced system monitoring.
- Engineered asynchronous task systems with Celery, improving task management for diverse corporate applications.
- Led integration of IoT, LoRa, WebSocket, and Itron platforms, delivering cutting-edge solutions for client utilities.
- Built secure Single Sign-On (SSO) features with OAuth2, JWT, Okta, and Microsoft to streamline user authentication.
- Developed OCR platform, processing pipelines utilizing PyTorch, Spacy, NLTK, GCP Document AI with Cloud Functions, AWS Textract with AWS Lambda, Tesseract, PaddleOCR and third party models (Hugging Face, OpenAI). Used Hugging Face, Spacy, NLTK for pre-processing, NLP parsing and error-detection. Implemented A/B testing for best outcomes.
- Delivered platform API MVP and new versions using .NET Core 7 and C#, modernizing back-end architecture, improving data layer and API service granularity.
- Provided the implementation of middleware, conversational messaging and notifications infrastructure in Rust and ZeroMQ for extended platform features.

Technologies: Python, Go, Angular, TypeScript, Rust, Flask, FastAPI, Django, PostgreSQL, RabbitMQ, Redis, Celery, WebSocket, D3.js, Three.js, HighCharts, Pandas, PyJanitor, OpenAPI, ELK Stack, Grafana, Datadog, Docker, Kubernetes, AWS (Textract, SageMaker), GCP, Terraform, Pulumi, Jenkins, Looker, Prometheus, PowerBI, Rust, ZeroMQ, .NET Core, C#, Azure Cosmos DB, Hugging Face, OpenAI, Okta, ZeroMQ, Dagster, Prefect, DBT, Pandas, Pydantic, Polars, Kafka, Debezium, LangChain, LlamaIndex, Kubeflow, Node.js.

1010data, New York, NY

- Architected and implemented ETL pipelines using Airflow DAGs, custom operators, and plugins ingesting data into AWS Redshift, Snowflake, Databricks, 1010data warehouse.
- Migrated Flask projects to Dockerized environments, leading orchestration, technology stack selection, and stakeholder negotiations.
- Designed and deployed event-driven infrastructures using Apache Kafka and AWS services for highly available solutions.
- Built and optimized HIPAA-compliant interoperability and FHIR/HL7/X12 data ingestion/integration and analysis, utilized SMART on FHIR to improve security and interoperability between medical data providers. Used Smile CDR for interoperability, data integration, CDS and automation.
- Built and optimized MLOps pipelines using PyTorch, Scikit-Learn, and Spacy, training and maintaining ML models.
- Automated infrastructure provisioning with Terraform and Ansible, ensuring scalability and reliability across multi-tier cloud architectures.
- Evaluated third-party tools, optimized existing architectures for cost efficiency, and ensured compliance with organizational policies.
- Collaborated cross-functionally to develop innovative solutions, meeting business goals while managing multiple concurrent projects. Provided guidance on compliance.
- Monitored and optimized system performance, implemented disaster recovery plans, and maintained comprehensive documentation of architectures and configurations.
- Consulted data engineering teams on best practices, stacks, tradeoffs, scalability, deployments on architectures for ETL pipelines using warehouses and data lakes.

Technologies: Apache Airflow, Luigi, Flask, Docker, Kubernetes, Terraform, Ansible, Kafka, AWS, PostgreSQL, CockroachDB, ELK Stack, Pandas, PyJanitor, NumPy, scikit-learn, Apache Spark, Scala, PySpark, Snowflake, BigQuery, DBT, Databricks, Vue, React, Neo4J, Tableau, Microsoft PowerBI, Python, Go, Java, Spring Boot, .NET Core, C#, Apache Solr, Lucene, Elasticsearch, Looker, Sentry, DVC, MLflow, Kubeflow, Dagster, Prefect, DBT, OpenLineage, HAPI FHIR, X12

Senior Developer

July 2016 – July 2018

Thirstie, New York, NY

- Built and scaled a marketplace platform using Tornado, Django, and Celery, with features like currency handling, parallel processing, and optimized concurrency (asyncio, multiprocessing).
- Designed and implemented the platform's data schema, models, and authorization service with JWT, including SSO evaluation.
- Conducted business analysis to prioritize features, align resources, and ensure timely delivery of technical tasks.
- Automated CI/CD pipelines and infrastructure management using Jenkins, Docker, Ansible, and Supervisor.
- Migrated legacy infrastructure, codebases, and schemas, ensuring seamless data transitions with Pewee, MySQL, and yoyo-migrations.
- Integrated third-party services including Braintree Payments, TaxJar, and Twilio to enhance platform functionality.
- Optimized system performance through data type tuning, schema normalization, and Redis enhancements.

Dmitry Roitman 150 West End Avenue, Brooklyn, NY 11235 | Phone: 718 404 6471 | dmitryro@gmail.com

LinkedIn: <https://www.linkedin.com/in/dmitry-r-4104b56/> | Website: <https://linktr.ee/dmitryro>

| GitHub: <https://github.com/dmitryro>

- Developed a gRPC-based integration layer and implemented load-time optimizations for scalability and reliability.

Technologies: Tornado, Django, Flask, MySQL, PostgreSQL, gRPC, AWS (EC2, S3, RDS), Digital Ocean, Elasticsearch, Redis, Docker, Celery, RQ, React, Apollo, GraphQL, Jenkins, Ansible, ELK Stack, New Relic.

ADDITIONAL EXPERIENCE:

Senior Developer / Architect

Art Revolution, New York, NY (Jan 2015 – Jul 2016)

- Built a Django-based social platform with advanced media processing, messaging, and integrations (OAuth2, UPS, USPS, Braintree).
- Developed AI/ML pipelines for image classification using Theano, Keras, and TensorFlow.
- Performed DevOps tasks including CI/CD pipelines, Docker, Ansible, and AWS infrastructure.

Integration Developer / Architect

Signetcs, New York, NY (Jan 2014 – Dec 2014)

- Designed and implemented B2B integration systems using Python, Node.js, and SOAP.
- Built custom job-handling and log analysis tools leveraging Redis, Tornado, and D3.
- Configured system security with Kerberos and OpenSSL.

Senior Developer

CPX Interactive, New York, NY (Sep 2013 – Dec 2013)

- Developed RESTful APIs for feed aggregation and integrated social APIs.
- Built scalable backend infrastructure with Django, Tornado, and Redis.

Senior Consultant

Northpoint Solutions, New York, NY (May 2013 – Aug 2013)

- Redesigned Rotary International's large-scale community site using Drupal.
- Implemented custom modules, group memberships, internationalization, and Google Maps integration.

Developer / Consultant

Analog Method, New York, NY (Feb 2013 – May 2013)

- Refactored Drupal-based sites with custom modules and interactive forms.
- Integrated SOAP-based e-commerce and geolocation functionality.

Independent Consultant / Freelancer

New York, NY (Feb 2009 – Dec 2014)

- Developed and maintained websites using WordPress, Drupal, and custom PHP/JavaScript solutions.
- Enhanced donation workflows and membership features for Covenant House (Drupal 6).

Java Software Developer

Algorithmic Creations, New York, NY (Jan 2001 – Feb 2009)

- Developed payment processing systems for ClickAndBuy with Java, J2EE and Hibernate, Spring Framework, JBoss and Apache Tomcat.
- Integrated third-party APIs (Axis 1.4) and optimized backend databases (MySQL, Oracle).